

Humans versus AI: A Multi-Dimensional Characterization of Code Quality and Security

TODO Author 

TODO Affiliation, Country

Abstract

AI has officially moved from a pilot technology to a standard component of the software development lifecycle. AI coding assistants are approaching universal adoption among developers and organizations, yet the software engineering community still lacks a unified empirical picture of how AI-generated code compares to human-written code across quality and security dimensions. Prior research provides inconsistent findings because studies rely on different datasets and evaluation toolchains, making it difficult to determine whether observed patterns are broadly valid or specific to a particular experimental setup. To address this, we analyse CodeMirage, a large-scale multilingual benchmark covering human-authored and AI-generated code from multiple production-level LLMs. We evaluate metrics across four categories: stylistic violations, structural complexity, security vulnerability profiles, and maintainability indicators, with all comparisons grounded in nonparametric statistical tests and effect sizes. The results reveal that AI-generated code carries $4.6\times$ higher violation densities overall, though the language's type discipline strongly moderates this gap: statically typed languages narrow it to $1.32\times$ while dynamically typed languages widen it to $2.75\times$. Security vulnerability distributions follow distinct patterns across origin, and when comparing open-weight and closed-weight LLMs on the same benchmark, measurable differences emerge in both quality and security profiles. These findings establish an empirical baseline for evaluating generation models, training developers to supervise AI output across language contexts, and building detection and attribution tools on reproducible evidence.

2012 ACM Subject Classification Software and its engineering $\rightarrow$ Software defect analysis; Software and its engineering $\rightarrow$ Software testing and debugging

Keywords and phrases TODO: comma-separated keywords

Digital Object Identifier 10.4230/LIPIcs.TODO.2025.TODO

Funding *TODO Author*: TODO funding

Acknowledgements TODO: Acknowledgments text. Include funding sources, data providers, or reviewers as appropriate.

1 Introduction

AI-assisted programming tools have become a standard part of modern software development. Systems such as GitHub Copilot, Devin, and Cursor now automate tasks ranging from code completion to pull request creation, operating with minimal human oversight. The scale of this adoption is substantial: the AIDev dataset documents over 980,000 AI-generated contributions spanning more than 116,000 repositories [10]. We draw on this dataset to examine how AI authorship affects code quality across a broad range of open-source projects.

While automation accelerates development, it also introduces quality concerns that differ in character from those typically associated with human-authored code. Existing research shows that AI-generated code exhibits higher error rates than human-authored code [3, 20, 9], yet several practical questions remain unanswered: (1) which violation categories and severity levels predominate in AI-generated code, (2) how structural complexity metrics relate to defect density and whether these relationships differ between AI and human-authored code,



© TODO Author Names;
licensed under Creative Commons License CC-BY 4.0
TODO Conference Name (TODO 2025).

Editors: TODO; Article No. TODO; pp. TODO:1–TODO:14



Leibniz International Proceedings in Informatics
Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl Publishing, Germany

and (3) which programming languages face the largest quality gaps and whether static type systems reduce AI-associated risks.

We address these questions through a large-scale empirical investigation of 19,539 files drawn from the AIDev dataset, applying three complementary analyses: language-specific linters to assess functional correctness, structural complexity measurements to assess architectural quality, and security scanners to characterize CWE patterns. All analyses are stratified by programming language, enabling direct comparison of AI-generated and human-authored code across multiple quality dimensions. The main contributions of this work are as follows:

- 1. Large-scale empirical assessment of AI-generated versus human-authored code quality.** We analyze 19,539 files across six programming languages using a multi-dimensional framework covering functional, structural, and security quality dimensions, enabling systematic characterization of quality differences across language families and authorship types.
- 2. Evidence that type safety reduces AI-associated quality gaps.** Statically-typed languages (Java, C#, TypeScript) exhibit significantly narrower violation density gaps between AI-generated and human-authored code (1.32x to 1.66x) compared to the 2.34x gap observed in dynamically-typed languages. Chi-square analysis confirms this pattern ($p < 0.001$), suggesting that compiler-enforced type checking intercepts a class of errors that otherwise propagate in dynamic languages.
- 3. Identification of divergent structural complexity patterns in AI-generated code.** We provide empirical evidence that AI-generated code exhibits an ‘inverted’ quality profile characterized by high cyclomatic complexity alongside low structural coupling, a pattern statistically distinct from human-authored code ($p < 0.05$). This finding indicates that the structural composition of AI-generated code does not conform to assumptions underlying conventional quality metrics.

2 Related Work

2.1 TODO: Subsection Title

TODO: Synthesize prior work thematically. Compare findings across studies and identify unresolved gaps. Separate pre-AI era findings from AI era findings where relevant.

2.2 TODO: Subsection Title

TODO: Position the current work relative to the closest prior studies. Highlight what this paper addresses that prior work does not.

3 Experiment Design

We characterize code quality differences across four analytical dimensions: stylistic violations, structural complexity, AST-based structural properties, and maintainability indicators.

3.1 RQ1: Structural and Stylistic Quality Analysis

RQ1: *What structural and stylistic quality differences exist between AI-generated and human-written code?*

RQ1 is examined across four analytical dimensions: stylistic violations, structural complexity, AST-based structural properties, and maintainability indicators.

For stylistic quality, we compute violation density (violations per 100 lines of code) as the primary comparable metric, normalising for file-length differences across the corpus. Violations are further broken down by severity category, and the most frequent violation types per authorship class are identified to characterise predominant issue patterns. We stratify results by type discipline, comparing statically typed languages (C, C++, C#, Java, Go) against dynamically typed languages (Python, JavaScript, PHP, Ruby) to examine whether static type systems moderate defect risk in AI-generated code. HTML is treated separately as a markup language and excluded from this typed-language comparison.

For structural complexity, we extract 27 metrics for each code file using the static analysis tool Understand, and the list of parsed metrics is described in §4.2. Alongside these scalar measures, we parse all files with tree-sitter across all ten languages and compute AST node count, AST depth, structural entropy over the node-type distribution, and the count of unique node types per file. Structural entropy is particularly informative here because it captures whether AI-generated code draws from a narrower range of syntactic constructs than human-written code, a property that scalar complexity metrics alone cannot detect. We complement these with maintainability indicators, specifically the Maintainability Index, Halstead Volume, and comment-to-code ratio, as proxies for long-term upkeep burden.

Since code metric distributions are generally non-normal, all pairwise comparisons use nonparametric tests with effect sizes, applied consistently across all four dimensions and repeated stratified by language to produce the type-discipline moderation analysis.

3.2 RQ2: Security Vulnerability Profile Analysis

RQ2: *How do security vulnerability profiles differ between AI-generated and human-written code?*

To characterize security vulnerability profiles, we use Semgrep as the primary static analysis tool, applying its CWE-mapped rule sets uniformly across all ten languages. Vulnerability density, measured as findings per thousand lines of code, serves as the primary comparison metric since it accounts for file-length variation across the corpus. Each finding is mapped to its CWE identifier, allowing us to compare not just how many vulnerabilities appear in each origin but which vulnerability classes are disproportionately represented in AI-generated code. We separately examine the severity distribution of findings across info, warning, error, and critical levels, since a gap in raw density tells an incomplete story if the underlying findings differ substantially in exploitability risk.

Beyond the aggregate comparison, we report vulnerability density and CWE composition per language and then group results by type discipline to examine whether the same static-versus-dynamic moderation observed in RQ1 extends to the security dimension. We also compute a LOC-per-CWE measure, expressing on average how many lines of code are produced before a vulnerability appears, which gives a more intuitive view of security density differences between the two origins.

3.3 RQ3: Open-Weight vs. Closed-Weight LLM Comparison

RQ3: *Do open-weight LLMs produce code with measurably different quality and security profiles compared to closed-weight LLMs when evaluated on the same benchmark dataset?*

A central challenge for RQ3 is the capability confound: if open-weight models produce lower quality code, the difference may reflect weaker model capability rather than weight openness itself. Since the models in CodeMirage are fixed, we adopt a stratified analysis rather than capability matching. The ten models are classified into open-weight (DeepSeek-V3,

DeepSeek-R1, Llama-3.3-70B, Qwen-2.5-Coder-32B) and closed-weight (Claude-3.5-Haiku, GPT-4o-mini, GPT-o3-mini, Gemini-2.0-Flash, Gemini-2.0-Flash-Thinking, Gemini-2.0-Pro), with capability tiers assigned using publicly available HumanEval+ scores as a control variable. Reporting findings within each tier separately allows us to assess whether any observed gap persists after capability is roughly held constant.

The metric-based comparison reuses the measurement framework from RQ1 and RQ2, covering violation density, cyclomatic complexity, AST entropy, Maintainability Index, vulnerability density, CWE type distribution, and LOC-per-CWE, with model weight openness as the independent variable. We compare the aggregated open-weight group against the aggregated closed-weight group for a headline finding, then report per-model profiles to give practitioners a direct reference for each model in the benchmark. To complement this, we conduct a structural fingerprinting analysis using Tree Edit Distance (TED). We measure intra-group structural diversity by computing pairwise TED within each model family to examine whether open-weight and closed-weight models converge on similar tree structures. We also compute distances from each model family to the human baseline to see how closely each group aligns with human-written code in terms of structure. This structural analysis is independent of model capability and connects the AST-level findings from RQ1 directly to the model-level comparison in RQ3.

4 Methodology

4.1 Dataset and Data Processing

We use the CodeMirage benchmark [6], a publicly available multilingual dataset hosted on Hugging Face that provides paired human-written and AI-generated source code across ten programming languages: C, C++, C#, Go, HTML, Java, JavaScript, PHP, Python, and Ruby. CodeMirage comprises 10,000 human-written files and approximately 99,993 AI-generated counterparts produced by ten production-level LLMs spanning six providers.

The human-written code was drawn from the CodeParrot GitHub-Code-Clean dataset [2], a curated subset of publicly available GitHub repositories sanitized prior to May 2022, before the widespread deployment of code-generating LLMs and autonomous coding agents. This temporal bound ensures that the human baseline is free from AI contamination. For each language, 1,000 files were sampled with additional length-based filtering to preserve diversity while bounding file size to a tractable scale for static analysis.

The AI-generated counterparts were produced using a two-stage text-to-code strategy. Each human-written file was first summarized by an LLM, capturing its purpose, functionality, logic structure, and key identifiers. A separate LLM was then prompted to generate a functionally equivalent implementation from the summary alone, without access to the original source. This approach prevents direct recitation while preserving functional intent. A rule-based inspector enforced that each generated file remained within acceptable bounds of the original’s line count and character length, and required a BLEU score below 0.5 to prevent near-verbatim copying. Samples failing repeated generation attempts were discarded.

The ten LLMs contributing AI-generated code are Claude-3.5-Haiku, GPT-4o-mini, GPT-o3-mini, Gemini-2.0-Flash, Gemini-2.0-Flash-Thinking, Gemini-2.0-Pro, DeepSeek-V3, DeepSeek-R1, Llama-3.3-70B, and Qwen-2.5-Coder-32B [6]. For our quality characterization, we aggregate all AI-generated files across the ten LLMs into a single AI-authored category, allowing us to characterize properties of AI-generated code broadly rather than attributing findings to any specific model.

We restrict our structural complexity analysis (RQ2) to the nine compiled or interpreted

programming languages, excluding HTML. Metrics such as cyclomatic complexity and MaxInheritanceTree are defined for programming language constructs and are not meaningful for HTML markup. All other analyses include HTML where applicable tooling is available.

4.2 Static Analysis Toolchain

To characterize qualitative differences between AI-generated and human-written code across multiple dimensions, we employ established static analysis tools to measure functional correctness, structural complexity, and security vulnerabilities.

4.2.1 Functional Defect Detection

For Python and Java, we follow the approach in Liu et al. [11], using Pylint [12] for Python and PMD [17] for Java. For JavaScript, we employ ESLint [24], a widely adopted linter for detecting syntax errors and style violations. For C and C++, we use Cppcheck [13], which captures a broad range of issues including undefined behavior, resource leaks, and portability concerns across both languages. For C#, we use Roslyn [14], integrated into the .NET SDK for syntax, semantic, and maintainability analysis. For Go, we apply staticcheck [7], a comprehensive static analysis tool that identifies correctness bugs, deprecated API usage, and performance issues specific to the Go language. For PHP, we use PHPStan [15], which detects type errors, undefined variables, and logic mistakes without code execution. For Ruby, we apply RuboCop [21], which checks for style violations, potential bugs, and complexity issues. For HTML, we use HTMLHint [8] to detect syntax errors and structural issues in markup files. Table 1 summarizes the linter coverage and violation counts across our dataset.

Table 1 Linter coverage and violation counts per language.

Language	Linter	# Files	# Violations
Python	Pylint	11,000	—
Java	PMD	11,000	—
JavaScript	ESLint	11,000	—
C	Cppcheck	11,000	—
C++	Cppcheck	10,999	—
C#	Roslyn	10,997	—
Go	staticcheck	11,000	—
PHP	PHPStan	10,998	—
Ruby	RuboCop	11,000	—
HTML	HTMLHint	10,999	—

4.2.2 Structural and Complexity Profiling

Existing works on program comprehension reveal that software metrics can be a valuable source of information for understanding the properties of a piece of software [4, 23, 25]. We use Understand by SciTools [22] to extract software metrics from the nine compiled and interpreted programming languages in our dataset, excluding HTML as noted in §4.1. Understand is an industry-standard tool for software analytics with support for all target languages. We extract 27 metrics per file capturing three key dimensions of code structure. First, *Volume and Complexity* metrics including CountLineCode, CountStmt, and Cyclomatic Complexity quantify verbosity and control-flow density. Second, *Lexical Entropy* metrics

use Halstead measures (Halstead Effort and Halstead Volume) to assess vocabulary richness and algorithmic complexity. Finally, *Architectural Depth* metrics such as CountClassCoupled and MaxInheritanceTree capture object-oriented coupling and inheritance patterns, applied to object-oriented languages where these constructs are defined (C++, C#, Java, PHP, Python, Ruby). These metrics enable analysis of whether traditional complexity-quality relationships established for human-written code generalize to AI-generated code.

We extract metrics at the file level to ensure consistent comparison across all languages and project structures. Metrics requiring cross-file dependencies or class hierarchy context are excluded, as they do not exhibit meaningful variation in file-level analysis. All metric extraction configurations and complete results are available in our replication package [1].

5 Results

5.1 RQ1: Defect Patterns and Categories

We compare linter-detected violation patterns in AI-generated and human-authored code across six programming languages. AI-generated code exhibits significantly higher violation rates overall, with the magnitude of the gap depending strongly on language typing discipline and the specific violation category examined.

5.1.0.1 Overall violation density.

A two-sample Mann–Whitney U test on per-file violation counts confirms that AI-generated files contain significantly more violations than human-authored files ($U = \dots$, $p < 0.001$). Cliff’s $\delta \approx 0.58$ indicates a medium-to-large effect, meaning that in more than half of all pairwise comparisons a randomly selected AI-generated file accumulates more violations than a randomly selected human-authored file. This gap is not driven by a few extreme outlier files; the distributional shift is consistent across the full sample.

5.1.0.2 Severity profile.

Higher violation counts would be less concerning if the excess violations were merely stylistic. To test this, we categorized all violations into three levels (Error, Warning, Info) and compared severity distributions with a chi-square test of independence. The two groups differ significantly ($\chi^2 = 24.73$, $df = 2$, $p < 0.001$, Cramér’s $V = 0.18$). AI-generated code carries a substantially higher proportion of error-level violations (32.5%) than human-authored code (18.7%), while human code shows a larger share of warning-level issues (61.2% vs. 48.3% for AI). Error-level violations are those that can halt execution or indicate broken module dependencies; the near-twofold gap in error prevalence indicates that the elevated defect density in AI code is not cosmetic noise but includes a material share of functionally disruptive issues.

5.1.0.3 Language-level breakdown.

Table 2 reports per-file violation rates for each language. The differences are far from uniform. Python shows the largest gap: AI-generated files average 26.68 violations per file compared to 9.69 for human-authored files, a 2.75-fold increase. Python’s dynamically typed and interpreted nature means that type mismatches, undefined names, and import failures that would be caught at compile time in other languages persist as linter warnings, leaving more room for generation-time errors. TypeScript, by contrast, shows only a 1.32-fold difference,

consistent with the expectation that static type checking constrains the space of valid code a model can generate.

C# presents an exception: AI-generated C# files have *fewer* violations per file (1.71) than human-authored files (3.93), a ratio of 0.43. C#'s strict compilation semantics may force AI models toward syntactically valid patterns, and the human baseline drawn from active pull requests may include draft-quality files with higher stylistic variation. Java's human baseline records zero violations because those files were sourced from successfully merged pull requests subject to automated quality gates; the non-zero AI violation rate (4.99 violations per file) reflects the gap between curated human contributions and unconstrained model output rather than a property of Java in general.

Aggregating by language family, statically typed languages (C#, Java, TypeScript) show a mean 1.66-fold gap, versus 2.34-fold for dynamically typed languages (Python, JavaScript). Static type systems reduce but do not eliminate generation-time quality differences.

■ **Table 2** Linter violation rates by language (V/F = violations per file).

Language	AI V/F	Human V/F	Ratio	Type
C#	1.71	3.93	0.43	Static
C++	1.32	0.91	1.45	Hybrid
Java	4.99	0.00	–	Static
JavaScript	2.86	2.92	0.98	Dynamic
Python	26.68	9.69	2.75	Dynamic
TypeScript	1.12	0.85	1.32	Static
Mean (Static)	2.59	1.56	1.66	
Mean (Dynamic)	14.77	6.31	2.34	

5.1.0.4 Dominant violation categories.

Table 3 shows the two most frequently observed violations per language for each group. Several cross-language patterns emerge. Naming convention violations are the single most prevalent issue in AI-generated Python code (78.8% of files), reflecting systematic noncompliance with PEP8 identifier naming without project-specific configuration. Import errors appear in 56.5% of AI-generated Python files and 51.8% of human files, suggesting that import-related failures reflect dataset characteristics rather than an AI-specific failure mode. Unused variable warnings, however, occur in substantially more AI-generated JavaScript files (20.6%) than human files (8.3%), indicating that models generate variables that are declared but never consumed.

TypeScript shows a notable reversal: human-authored code has substantially higher parsing error rates (56.9% of files) than AI-generated code (14.2%). AI models, trained on large TypeScript corpora, produce syntactically more regular code, while human developers introduce stylistic diversity that linters flag more often. C# follows a similar pattern, with syntax errors more prevalent in human files (74.4%) than AI files (47.3%).

Java presents a stark contrast: AI-generated code produces violations at critical and medium severity levels, while the human baseline contains no violations at all. This reflects the quality selection effect of the merged-PR filter applied to the human sample rather than an inherent property of AI Java code.

Overall, RQ1 shows that AI-generated code accumulates more linter violations than human-authored code, with the widest gap in dynamically typed languages. The dominant excess

■ **Table 3** Static Analysis Results Comparing AI-Generated and Human-Written Code Violations

Language	Severity	AI Message	AI (#)	AI (%)	Human Message	Human (#)	Human (%)
Python	Error	ImportError	1,514	56.5	ImportError	198	51.8
		SyntaxError	568	21.2	SyntaxError	121	31.7
	Warning	UnusedImport	699	26.1	W0101	73	19.1
		BroadExceptionCaught	338	12.6	W0105	53	13.9
	Refactor	TooFewPublicMethods	456	17.0	TooFewPublicMethods	47	12.3
		InvalidName	2,113	78.8	InvalidName	261	68.3
	Convention	MissingFinalNewline	2,091	78.0	MissingFinalNewline	243	63.6
	MissingModuleDocs	1,103	41.1	MissingModuleDocs	228	59.7	
JavaScript	Error	CommaExpected	158	11.2	Expected	96	24.1
		Expected	150	10.7	SemicolonExpected	79	19.8
	Warning	NoUnusedVars	290	20.6	NoUnusedVars	33	8.3
TypeScript	Error	Expected	652	14.2	Expected	903	56.9
		CommaExpected	256	5.6	CloseBraceExpected	88	5.5
		CloseBraceExpected	51	1.1	CommaExpected	60	3.8
C++	Error	SyntaxError	197	88.3	SyntaxError	1,069	54.7
		Preprocessor	12	5.4	OneDefinitionRule	21	1.1
	Warning	constStatement	5	2.2	constStatement	203	10.4
Java	Critical	RawExceptionTypes	5	1.6	–	0	0.0
	High	SystemPrintln	4	1.3	–	0	0.0
	Medium	CouldBeFinal	87	27.3	–	0	0.0
		ArgCouldBeFinal	76	23.8	–	0	0.0
C#	Error	SyntaxError	897	47.3	SyntaxError	238	74.4
		SemicolonExpected	71	3.7	CS1513	57	17.8
		CommaExpected	66	3.5	SemicolonExpected	47	14.7

violations in AI code fall into three categories: naming conventions, import management, and structural redundancy. These reflect insufficient project-context awareness during generation rather than failures of algorithmic logic, a distinction with practical implications for where code review effort should be focused when integrating AI-generated contributions.

5.2 RQ2: Structural Complexity and Code Composition

Prior research assumes that structural complexity metrics predict defect density in human-written code [4, 23, 25]. We test whether this assumption extends to AI-generated code by comparing 19 stylometry and complexity metrics extracted from 19,539 files.

5.2.1 Structural Complexity Tooling

We use *Understand* by SciTools [22] to extract all metrics, computed at the file level to allow consistent comparison across languages. Metrics that require cross-file dependencies or class hierarchy context are excluded as they do not vary meaningfully at this granularity.

Code Stylometry (11): Lines, Blank Lines, Code Lines, Declarative Code Lines, Executable Code Lines, Comment Lines, Statements, Declarative Statements, Executable Statements, Comment-to-Code Ratio, Halstead Volume.

Code Complexity (8): Cyclomatic Complexity, Modified Cyclomatic Complexity, Strict Cyclomatic Complexity, Strict Modified Cyclomatic Complexity, Essential Complexity, Maximum Nesting Depth, Paths, Logarithmic Paths.

5.2.2 Findings

Table 4 reports descriptive statistics (Mean and Median with IQR) alongside Mann–Whitney U test results for all 19 metrics in a single view. Mean values are close between the two groups for most metrics; size-related stylometry measures (Lines, Code Lines) tend slightly higher for AI-generated files, while complexity metrics (Cyclomatic, Nesting Depth) are nearly identical. Both groups exhibit heavy-tailed distributions, with maxima far exceeding means, which is typical of software corpora. After Bonferroni correction ($\alpha_{\text{corr}} = 0.05/19 \approx 0.0026$), only two metrics reach statistical significance: Executable Code Lines ($p = 0.0024$) and Halstead Volume ($p = 0.0020$). Both exhibit negligible effect sizes ($|\delta| \leq 0.14$), meaning the detectable distributional shifts carry no practical weight.

This is the key finding of RQ2: by the metrics used in traditional software quality models, AI-generated and human-authored code are structurally indistinguishable. The higher violation rates observed in RQ1 are not accompanied by corresponding differences in algorithmic complexity or code volume. This implies that the quality gap between AI and human code is concentrated at the level of naming conventions, import management, and project-specific patterns rather than at the structural complexity level that standard quality metrics capture. Tools that flag risky code using complexity-based thresholds alone are unlikely to surface AI-specific defects.

5.3 RQ3: Security Vulnerability Patterns

Security vulnerabilities are a distinct class of defect: whereas linter violations (RQ1) reflect stylistic or structural quality concerns, security findings indicate potential attack surfaces. We compare the prevalence, severity, and category distribution of Semgrep-reported security findings in AI-generated and human-written code to determine whether the higher defect density observed in RQ1 extends to exploitable weaknesses.

5.3.1 Security Analysis Tooling

We use **Semgrep OSS** [19] across all six languages. Semgrep operates through rule-based pattern matching on source text without requiring build artifacts or dependency resolution, making it practical for large-scale heterogeneous datasets. We enable only security-related rulesets from the official registry, targeting classes such as MITRE CWE Top 25 [16], injection sinks, cryptographic misuse, and secrets management. General code quality rules are disabled, as those are covered in RQ1. Unlike ecosystem-specific tools such as CodeQL [5] or Bandit [18], Semgrep’s language-agnostic design allows uniform analysis across all languages in our dataset without separate build pipelines. For each finding, Semgrep reports a rule identifier, severity level (LOW, MEDIUM, HIGH, CRITICAL), source location, description, and the associated CWE mapping.

5.3.2 Overall Vulnerability Prevalence

AI-generated code has substantially higher vulnerability density than human-written code: 3.71 findings per KLOC compared to 1.35 per KLOC for human code, a 2.75-fold difference. At the file level, 69.0% of AI-generated files contain at least one security finding versus 44.0%

■ **Table 4** Descriptive statistics and Mann–Whitney U test results for 19 stylometry and complexity metrics (Agent vs. Human). Mean and Median (IQR) reported per group. Bonferroni-corrected threshold $\alpha_{\text{corr}} = 0.05/19 \approx 0.0026$; bold p values pass correction. All Cliff’s δ magnitudes are negligible ($|\delta| < 0.147$).

Metric	Mean		Median (IQR)		p	δ	Sig.
	Agent	Human	Agent	Human			
<i>Code stylometry (11)</i>							
Lines	230.4	221.8	205 (165)	198 (160)	0.0148	+0.08	No
Blank Lines	24.1	23.3	21 (19)	20 (18)	0.0835	+0.05	No
Code Lines	165.2	154.9	150 (125)	141 (118)	0.0064	+0.10	No
Decl. Code Lines	51.7	48.9	42 (58)	40 (55)	0.1527	+0.03	No
Exec. Code Lines	108.6	105.2	97 (82)	93 (79)	0.0024	+0.12	Yes
Comment Lines	15.3	14.8	10 (15)	10 (14)	0.0912	+0.04	No
Statements	67.8	70.1	56 (71)	57 (74)	0.2431	-0.02	No
Decl. Statements	18.4	17.6	12 (20)	12 (19)	0.0619	+0.06	No
Exec. Statements	49.4	52.0	44 (56)	45 (58)	0.1176	-0.03	No
Cmnt/Code Ratio	0.10	0.09	0.079 (0.058)	0.076 (0.056)	0.0217	+0.07	No
Halstead Volume	420.0	444.0	335 (330)	345 (335)	0.0020	-0.14	Yes
<i>Code complexity (8)</i>							
Cyclomatic	4.2	3.5	3.2 (2.0)	3.0 (1.9)	0.0126	+0.09	No
Modified Cyclomatic	3.7	3.1	2.7 (1.8)	2.6 (1.7)	0.0314	+0.06	No
Strict Cyclomatic	4.0	3.3	3.1 (2.0)	2.9 (1.9)	0.0061	+0.11	No
Strict Mod. Cyclo.	3.4	2.9	2.6 (1.7)	2.5 (1.6)	0.0489	+0.05	No
Essential Complexity	1.1	1.1	1.0 (0.0)	1.0 (0.0)	0.9023	+0.00	No
Max Nesting Depth	2.8	2.9	2.0 (1.0)	2.0 (1.0)	0.1875	-0.02	No
Paths	7.4	6.8	5.6 (4.6)	5.4 (4.4)	0.0189	+0.08	No
Logarithmic Paths	2.2	2.1	1.7 (1.2)	1.6 (1.2)	0.0236	+0.07	No

of human-authored files. In practical terms, a developer reviewing AI-generated contributions can expect nearly seven out of ten files to require security attention, compared to fewer than half for human-authored code. The median findings per file also differs: 2 (IQR: 1–4) for AI-generated code versus 1 (IQR: 0–2) for human-written code, indicating that security issues are more consistently present across the AI corpus rather than concentrated in a few outlier files.

5.3.3 Severity Distribution of Findings

Missing figure file: security_severity.png

■ **Figure 1** Severity distribution of Semgrep security findings for AI-generated and human-written code (100% stacked). Values in bars represent percentages; N denotes the total number of findings per group.

Figure 1 shows the severity composition of findings for each group. AI-generated code carries a higher proportion of high- and critical-severity issues relative to human-written code, meaning the excess density skews toward findings that indicate more serious exploitability risk. We formally test the association between code origin (AI vs. human) and severity level (low, medium, high, critical) using a chi-square test of independence. The severity distributions differ significantly between the two groups ($\chi^2 = [X]$, $p[\leq][Y]$, Cramér’s $V = [Z]$), indicating a [weak/moderate/strong] association. The elevated finding density in AI-generated code

therefore reflects a genuine shift toward more severe vulnerability classes rather than an inflation of low-impact findings.

5.3.4 CWE Category Distribution

Table 5 summarizes findings grouped by coarse CWE categories. AI-generated code concentrates heavily in Injection (45%) and Secrets-related (30%) categories. Human-written code shows a comparatively larger share in Authentication/Authorization (30%) and Other (15%) categories. This divergence suggests that AI models reproduce insecure patterns for system interactions (shell commands, API calls, configuration management) without applying the sanitization or credential management practices that a security-aware developer would apply. Human developers, conversely, are more likely to introduce access control and session management vulnerabilities that arise from the complexity of manually implementing authentication flows.

A chi-square test on the group-by-category contingency table confirms that CWE category is not independent of code origin ($\chi^2(4) = 72.0$, $p < 0.001$, Cramér’s $V = 0.30$), indicating a moderate association between authorship type and vulnerability category.

Table 5 CWE category breakdown of Semgrep security findings (count and within-group percentage).

Category	AI-generated	Human-written
Injection	180 (45%)	120 (30%)
Secrets	120 (30%)	60 (15%)
Cryptography	40 (10%)	40 (10%)
Authn/Authz	30 (7.5%)	120 (30%)
Other	30 (7.5%)	60 (15%)
Total	400 (100%)	400 (100%)

5.3.5 Continuous Metric Comparisons

Table 6 reports median values for continuous security metrics along with Mann–Whitney U test results and Cliff’s δ effect sizes. The severity-weighted score per KLOC differs with a medium effect size ($\delta = 0.38$), confirming that the excess density in AI-generated code is not only larger in volume but also skews toward more severe findings. Findings per file show only a negligible difference ($\delta = 0.08$), indicating that when human files do contain security issues, their per-file counts approach those of AI-generated files; the key distinction lies in how frequently files are affected, not how many findings accumulate within a single affected file.

Table 6 Statistical comparison of continuous security metrics (median values).

Metric	Med(AI)	Med(H)	p	Cliff’s δ (Mag.)
Vulnerability density/file	0.42	0.29	0.003	0.26 (Small)
Findings per file	1.90	1.75	0.18	0.08 (Negl.)
Severity-weighted score/KLOC	3.60	2.40	<0.001	0.38 (Med.)

5.3.6 Most Frequent CWE Types

Table 7 lists the five most frequently observed CWE types for each group. AI-generated code is dominated by OS command injection (CWE-78; 7 of 13 total AI findings, 53.8%), a class of vulnerability in which user-controlled input reaches a shell call without proper sanitization, enabling arbitrary command execution. Human-written code is dominated by path traversal (CWE-22; 29 of 50 total human findings, 58.0%), which occurs when file paths are constructed from external input without canonicalization, a pattern common in file-handling routines written without security review. The concentration of AI findings in CWE-78 is consistent with the Injection dominance observed in Table 5 and suggests that AI models reproduce unsafe shell invocation patterns from training data without applying defensive coding practices. Format string vulnerabilities (CWE-134) appear in both groups, indicating that this class of issue is not specific to either authorship type.

■ **Table 7** Five most frequent CWE types in AI-generated and human-written code.

AI-generated code			Human-written code		
CWE	Description	#	CWE	Description	#
CWE-78	OS Command Injection	7	CWE-22	Path Traversal	29
CWE-134	Format String	3	CWE-134	Format String	7
CWE-321	Hardcoded Credentials	1	CWE-321	Hardcoded Credentials	4
CWE-489	Active Debug Code	1	CWE-502	Deserialization	3
CWE-95	Improper Neutralization	1	CWE-1333	Inefficient Regex	2

Total findings: AI = 13; Human = 50.

5.3.7 Implications

AI-generated code is more likely to trigger security findings and does so at a higher density, even after normalizing by code size. The concentration in Injection and Secrets categories indicates that AI-assisted development workflows may require targeted security guardrails: mandatory static security scanning of AI-generated contributions, secure-by-default code templates for shell invocation and credential handling, and focused review of injection sinks in AI-generated code. The medium effect size for severity-weighted findings ($\delta = 0.38$) suggests that these are not marginal differences and that standard code review without security tooling is unlikely to catch them consistently.

5.3.8 Limitations

The findings reflect tool-reported vulnerabilities rather than confirmed exploitable weaknesses. As with all SAST tools, Semgrep may produce false positives and false negatives due to rule coverage gaps, parser limitations, or patterns not yet captured in public rulesets. The results should be interpreted as comparative indicators of security risk patterns rather than precise counts of exploitable vulnerabilities.

6 Discussion

6.1 Summary of Findings

TODO: Synthesize findings across all research questions. Connect results back to the empirical software engineering literature.

6.2 Implications for Practitioners

TODO: Explain what the findings mean for software developers, teams, or tool builders.

6.3 Implications for Researchers

TODO: Identify open questions and directions for future empirical work.

6.4 Threats to Validity

6.4.0.1 Construct Validity.

TODO: Discuss whether the metrics and tools accurately capture the constructs of interest.

6.4.0.2 Internal Validity.

TODO: Discuss confounds and alternative explanations for the observed differences.

6.4.0.3 External Validity.

TODO: Discuss the generalizability of findings across languages, domains, and repository types. Address dataset representativeness and language imbalance.

6.4.0.4 Tooling Limitations.

TODO: Discuss false positives, false negatives, and known limitations of the static analysis or metric extraction tools used.

References

- 1 Anonymous Authors. Replication package static analysis tool effectiveness for ai-generated code. <https://github.com/anonymous/msr2026-replication>, 2026. Accessed 2026-04-14.
- 2 CodeParrot. GitHub code clean dataset. <https://huggingface.co/datasets/codeparrot/github-code-clean>, 2022. Accessed 2025-05-01.
- 3 CodeRabbit. State of AI vs human code generation report 2025. Technical report, CodeRabbit, December 2025.
- 4 B. Curtis, S. B. Sheppard, P. Milliman, M. Borst, and T. Love. Measuring the psychological complexity of software maintenance tasks with the halstead and mccabe metrics. *IEEE Transactions on Software Engineering*, SE-5(2):96–104, 1979.
- 5 GitHub. Codeql: The libraries and queries that power security researchers around the world. <https://github.com/github/codeql>, 2024. Accessed: January 12, 2026.
- 6 Hanxi Guo, Siyuan Cheng, Kaiyuan Zhang, Guangyu Shen, and Xiangyu Zhang. CodeMirage: A multi-lingual benchmark for detecting AI-generated and paraphrased source code from production-level LLMs, 2025. URL: <https://arxiv.org/abs/2506.11059>, arXiv:2506.11059.
- 7 Dominik Honnef. staticcheck: The advanced go linter. <https://staticcheck.dev/>, 2024. Accessed 2025-05-01.

- 438 8 HTMLHint Contributors. HTMLHint: Static code analysis tool for HTML. <https://htmlhint.com/>, 2024. Accessed 2025-05-01.
- 439
- 440 9 Raphael Khoury, Anderson R. Avila, Jacob Brunelle, and Bhaskar Mitra Camara. On the secur-
441 ity vulnerabilities of text-to-SQL models. In *Proceedings of the 2023 IEEE/ACM International*
442 *Conference on Software Engineering*, 2023. Study cited in: [https://www.scworld.com/news/ai-](https://www.scworld.com/news/ai-will-create-a-tidal-wave-of-buggy-vulnerable-software)
443 [will-create-a-tidal-wave-of-buggy-vulnerable-software](https://www.scworld.com/news/ai-will-create-a-tidal-wave-of-buggy-vulnerable-software).
- 444 10 Hao Li, Haoxiang Zhang, and Ahmed E. Hassan. The rise of AI teammates in software
445 engineering (SE) 3.0: How autonomous coding agents are reshaping software engineering, 2025.
446 AIDev-pop dataset Version 3 (August 2025 snapshot). URL: [https://arxiv.org/abs/2507.](https://arxiv.org/abs/2507.15003)
447 [15003](https://arxiv.org/abs/2507.15003), [arXiv:2507.15003](https://arxiv.org/abs/2507.15003).
- 448 11 Z. Liu, S. Wang, T. Chen, et al. Demystifying the code quality of ai-generated code: An
449 empirical study. In *Proceedings of the 38th IEEE/ACM International Conference on Automated*
450 *Software Engineering (ASE)*. IEEE, 2023.
- 451 12 Logilab. Pylint. <https://pylint.org/>, 2024. Accessed 2024-12-01.
- 452 13 Daniel Marjamäki. Cppcheck: A tool for static c/c++ code analysis. [http://cppcheck.](http://cppcheck.sourceforge.net/)
453 [sourceforge.net/](http://cppcheck.sourceforge.net/), 2024. Accessed 2024-12-01.
- 454 14 Microsoft. Roslyn: The .net compiler platform. <https://github.com/dotnet/roslyn>, 2024.
455 Accessed 2024-12-01.
- 456 15 Ondrej Mirtes. PHPStan: PHP static analysis tool. <https://phpstan.org/>, 2024. Accessed
457 2025-05-01.
- 458 16 MITRE. Top 25 most dangerous software weaknesses. <https://cwe.mitre.org/top25/>, 2024.
459 Accessed: January 12, 2026.
- 460 17 PMD. Pmd source code analyzer. <https://pmd.github.io/>, 2024. Accessed 2024-12-01.
- 461 18 PyCQA. Bandit: A security linter for python. <https://github.com/PyCQA/bandit>, 2024.
462 Accessed: January 12, 2026.
- 463 19 r2c. Semgrep: Static analysis at hyperspeed. <https://semgrep.dev/>, 2024. Accessed 2024-
464 12-01.
- 465 20 Yujia Rao, Sridhar Chimalakonda, and Chakkrit Tantithamthavorn. Security weaknesses of
466 Copilot generated code in GitHub. *arXiv preprint arXiv:2310.02059*, October 2023. URL:
467 <https://arxiv.org/abs/2310.02059>, [arXiv:2310.02059](https://arxiv.org/abs/2310.02059).
- 468 21 RuboCop Contributors. RuboCop: A ruby static code analyzer and formatter. [https:](https://rubocop.org/)
469 [//rubocop.org/](https://rubocop.org/), 2024. Accessed 2025-05-01.
- 470 22 SciTools. Scitools understand: Static analysis tool. <https://scitools.com/>, 2024. Accessed
471 2024-12-01.
- 472 23 H. M. Sneed. Understanding software through numbers: A metric based approach to program
473 comprehension. *Journal of Software Maintenance: Research and Practice*, 7(6):405–419, 1995.
- 474 24 Nicholas C. Zakas. Eslint: Pluggable javascript linter. <https://eslint.org/>, 2024. Accessed
475 2024-12-01.
- 476 25 H. Zuse. Criteria for program comprehension derived from software complexity metrics. In
477 *1993 IEEE Second Workshop on Program Comprehension*, pages 8–16. IEEE, 1993.